

# **C++ Course 12 : Move semantics.**

**2020 by Oleksiy Grechnyev**

# Copy operations (C++ 98), Move operations (C++ 11+)

```
String a("Mickey mouse");
```

```
String b(a);    // Copy Ctor
```

```
String c;
```

```
String c = a;    // Copy Assignment
```

```
b = string("Cinderella"); // Copy assignment in C++ 98, Move assignment in C++ 11+
```

```
c = move(b);    // Move assignment in C++ 11+, b is now empty string !
```

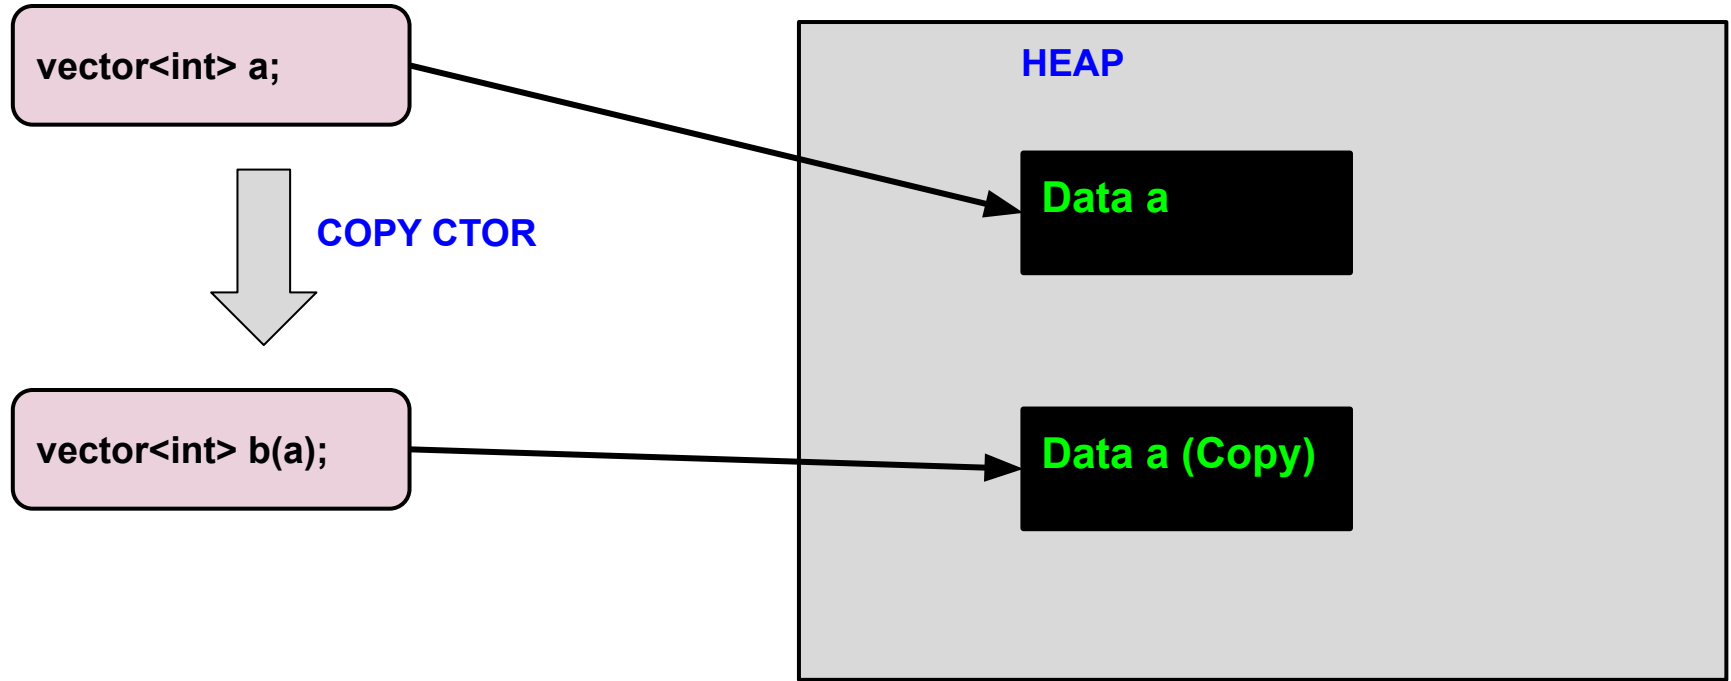
```
String d(move(a)); // Move ctor in C++ 11+, a is now empty string !
```

Copy is easy to understand: we clone our object.

But what is MOVE ?


DATA is moved, NOT OBJECT !!!

# Copy constructor : `std::vector` copies data in the heap




A new vector **b** is created. The data of vector **a** is COPIED in the heap to **b**.  
Copying class objects is often EXPENSIVE.

# Move constructor : data is moved to another vector object



# Move constructor : data is moved to another vector object

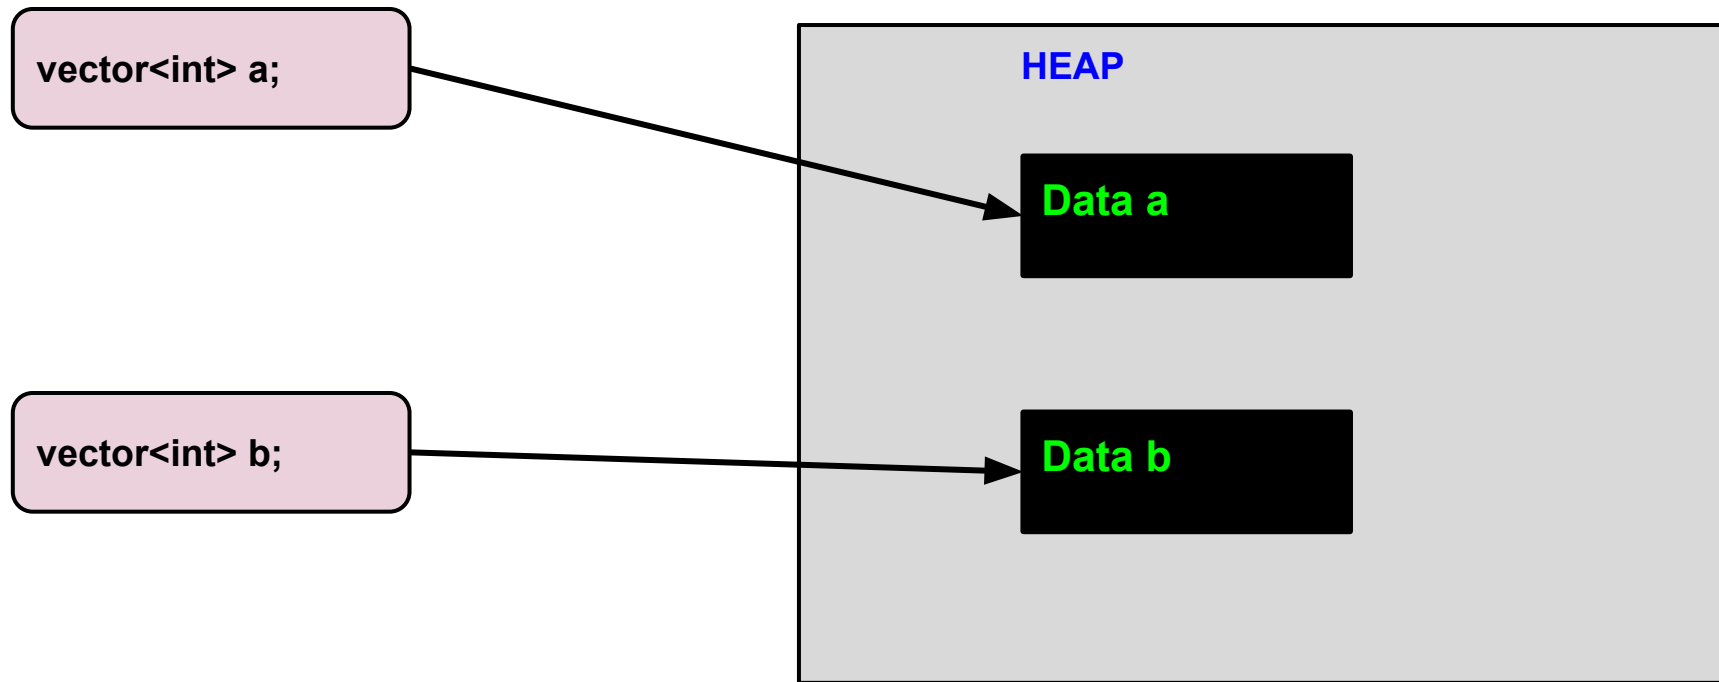


The data is moved from object **a** to the new vector **b**.

Vector **a** is now empty (**size == 0**, **capacity == 0**), but NOT destroyed !

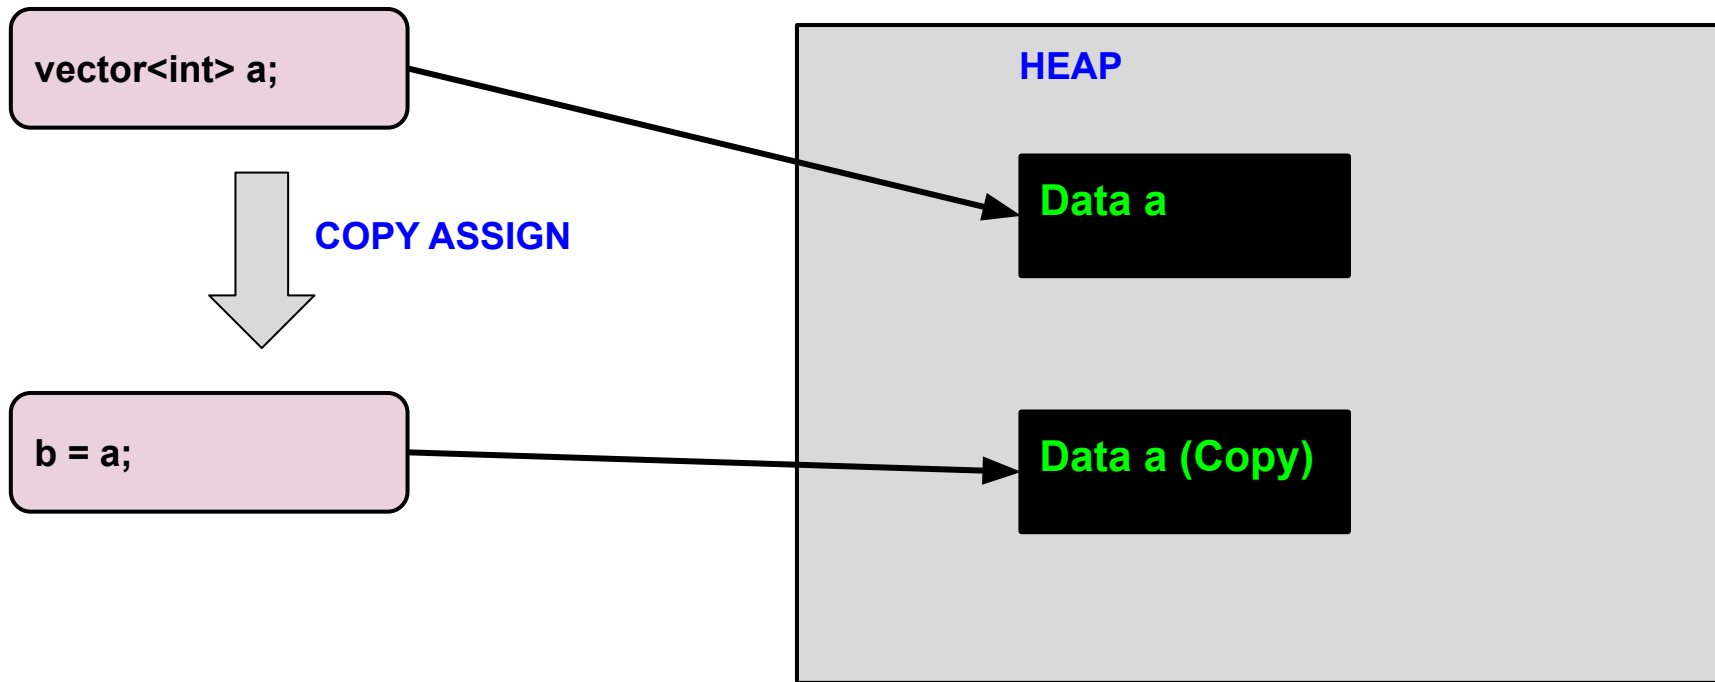
Data is moved, not objects !

# Before Copy or Move assignment



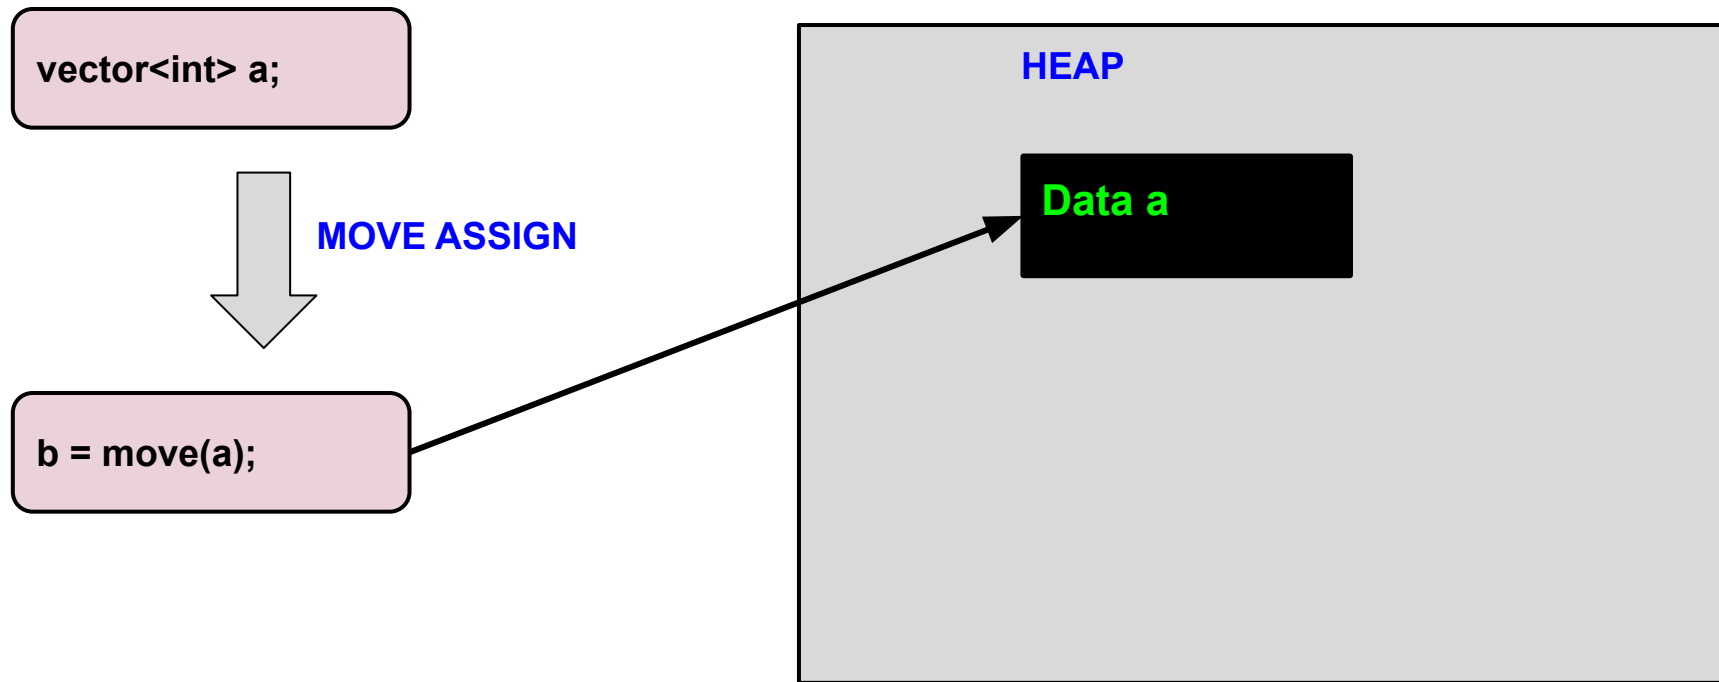
Vectors **a** and **b** both exist and keep their own heap data.

# Copy assignment



Data **b** is lost (and destroyed on the heap), Data **a** is copied in the heap to vector **b**.

# Move assignment



Data **b** is lost (and destroyed on the heap), Data **a** is moved to vector **b**.  
Vector **a** is now empty (**size == 0**, **capacity == 0**), but NOT destroyed !  
Data is moved, not objects !



# In a move operation the victim is NOT destroyed !!!

Move is a robbery, not murder !

```
{  
    string a("Sword");  
    string b(move(a));  
    // Now b == "Sword", a == "" (Empty, Not destroyed !)  
    cout << "a = " << a << endl; // OK, Empty  
    cout << "b = " << b << endl; // OK, Sword  
    a = "Spear"; // a can be assigned again  
} // Now a, b run out of scope and are destroyed !
```

Move victim is never destroyed only robbed of data (becomes empty/null).

Data is moved, not objects !

When is MOVE useful?

1. Classes with data in heap : **vector**, **string**, most other containers.
2. Classes which cannot be copied : **unique\_ptr**, **ifstream**, **thread**, **future**.

# Rvalue vs lvalue

Move operations are implemented with the help of RVALUE REFERENCES.

The terminology LVALUE, RVALUE comes from the early days of C.

The assignment statement:

**LVALUE = RVALUE;**

LVALUE (left value) can be assigned. Variable, array element **arr[i]**, etc.

RVALUE (right value) cannot be assigned. Temporary: literal, expression result etc.

However, in modern sense of the word:

**const a = 17;**

**a** is an LVALUE (every variable is LVALUE), but cannot be assigned!

Confusing!

We need a better definition of LVALUE, RVALUE !

BOOK: Scott Meyers, Effective Modern C++, chapter 5

# Rvalue vs lvalue: Better definition

LVALUE is something you can take the address of with the **&** operator:

```
string s("Mouse");
```

```
string * pS1 = & s; // OK, s in an LVALUE
```

```
string * pS2 = & string("Rat"); // ERROR, string("Rat") in an RVALUE
```

LVALUE examples:

1. Variable, or constant variable: **int a = 17;**
2. Function/method return by LVALUE ref : **arr.at(j)**
3. Operator return by LVALUE ref : **arr[i], \*ptr**

RVALUE examples:

1. Literal : **17, "Mouse"**
2. Temporary object : **string("Rat")**
3. Expression result: **2\*a + b**
4. Function return by value : **cos(alpha)**

# LVALUE vs RVALUE references

C++ 98 : Non-const (LVALUE) reference binds to LVALUE only.

```
string s("Guinea Pig");
```

```
string &lrS1 = s; // OK, ref to variable s
```

```
string &lrS2 = string("Marmot"); // ERROR !!!
```

C++ 98 : Const (LVALUE) reference binds to LVALUE, but also RVALUE.

```
string s("Guinea Pig");
```

```
const string &clrS1 = s; // OK, const ref to variable s
```

```
const string &clrS2 = string("Marmot"); // OK, const ref to a temporary
```

C++ 11+ : RVALUE reference binds to RVALUE (temporary object) only.

```
string s("Guinea Pig");
```

```
string &&rrS1 = s; // ERROR rvalue ref to lvalue !
```

```
string &&rrS2 = string("Marmot"); // OK, rvalue ref to a temporary
```

Note: The lifetime of TEMPORARY is extended until } for **clrS2** and **rrS2** !

# Copy and move constructors and assignment operators

```
Tjej(const Tjej & rhs) noexcept : name(rhs.name) {}           // Copy Ctor

Tjej(Tjej && rhs) noexcept : name(std::move(rhs.name)) {}     // Move Ctor

Tjej & operator= (const Tjej & rhs) noexcept{                // Copy assignment
    if (this != &rhs)    // Check for self-assignment
        name = rhs.name;
    return *this;
}

Tjej & operator= (Tjej && rhs) noexcept{                      // Move assignment
    if (this != &rhs)    // Check for self-assignment
        name = std::move(rhs.name);
    return *this;
}
```

**Tjej &&** is an *RVALUE* reference. **const Tjej &** is a const *LVALUE* reference.

The actual MOVE operation is implemented in move Ctor and Assignment.

**move()** does not move anything, it selects the *RVALUE* overload.

# Copy (LVALUES) and Move (RVALUES)

How are COPY or MOVE overloads selected ?

**String a("Mickey mouse");**

**String b(a);**        // Copy Ctor, a is a LVALUE

**String c;**

**String c = a;**        // Copy Assignment, a is a LVALUE

**b = string("Cinderella");**    // Move assignment, **string("Cinderella")** is a RVALUE

**c = static\_cast<string &&>(b);**        // Move assignment, cast to RVALUE !

**c = move(b);**        // The same

**String d(static\_cast<string &&>(a));**    // Move ctor, cast to RVALUE !

**String d(move(a));**        // The same

**move()** does not move anything, it casts to *RVALUE* !

Copy Ctor/assignment is selected for LVALUES.

Move Ctor/assignment is selected for RVALUES.

# Source (factory) : creates a new object

Return by value, no MOVE !

```
Tjej source1(const string & s){  
    Tjej t(s);  
    return t;  
}
```

...

```
Tjej t1 = source1("Karen Koenig");
```

Return Value Optimization : No copy or move (but not guaranteed until C++ 17) !

We can even make source within source:

```
Tjej source2(const string & s){  
    return source1(s);  
}
```

```
Tjej t2 = source2("Alice Elliot");
```

Still no copy or move, only 1 object created !

# Is RVALUE reference VARIABLE an RVALUE? NO !!!

```
void doSomething(Tjej && rrT){    // rrS is an RVALUE parameter
    string && rrS = string("Maria Traydor");    // rrS is an RVALUE variable
    // rrT and rrS are RVALUE refs, but not RVALUES !
    Tjej t(rrT);    // Copy, NOT MOVE !
    string s(rrS);    // Copy, NOT MOVE !
    ...
}
```

All variables and parameters are LVALUES by definition !

If we want move, we must cast them to rvalue refs as usual :

```
void doSomething2(Tjej && rrT){
    string && rrS = string("Maria Traydor");
    Tjej t(move(rrT));    // MOVE !
    string s(move(rrS));    // MOVE !
    ...
}
```



# Writing a sink (consumer)

A sink consumes an object by MOVE and lets it die.

```
void sink(Tjej && t0) {      // By rvalue ref
    Tjej t(move(t0));       // Move to a local var
    cout << "sink : " << t.getName() << endl;
} // t dies here
```

A copy of a ref parameter? Better to pass by value !

```
void sink1(Tjej t) {       // By value (move), local that dies !
    cout << "sink1 : " << t.getName() << endl;
} // t dies here
```

...

```
Tjej t3("Lucia");
sink1(move(t3)); // 1 move, 0 copy, Move ctor is selected for the parameter t
```

# Writing a sink (consumer)

A copy of a ref parameter? Better to pass by value !

```
void sink1(Tjej t) {           // By value (move), local that dies !  
    cout << "sink1 : " << t.getName() << endl;  
} // t dies here
```

...

```
Tjej t3("Lucia");  
sink1(move(t3)); // 1 move, 0 copy, Move ctor is selected for the parameter t
```

Chain sinks :

```
void sink2(Tjej t) { // By value (move), local that dies !  
    sink1(move(t)); // The correct way, move to die !  
} // t dies here
```

...

```
Tjej t4("Margarete Gertrude Zelle");  
sink2(move(t4)); // 2 moves, 0 copy, Move ctor is selected for the parameter t
```

# In Move Ctor/assign use move() for :

Class fields:

```
Tjej(Tjej && rhs) noexcept : name(std::move(rhs.name)){}    // Move Ctor
Tjej & operator= (Tjej && rhs) noexcept{                    // Move assignment
    if (this != &rhs)    // Check for self-assignment
        name= std::move(rhs.name);
    return *this;
}
```

Parent class:

```
class LilTjej : public Tjej{
    ...
    // Move Ctor
    LilTjej(LilTjej && rhs) noexcept : Tjej(std::move(rhs)){}
};
```

Remember ! Move ctor must be **noexcept** to work in **std::vector** !

# Move summary

1. MOVE operations are only useful for classes with heap memory (containers) and the ones which cannot be copied (**unique\_ptr**, **thread**, **ostream**).
2. C++ does not move anything, the move ctor/operator= does!
3. Move ctor/assign is selected for RVALUES, copy ctor/assign is selected for LVALUES.
4. **std::move()** does not move anything, it casts to RVALUE.
5. Sources return by value, sinks get parameter by value.
6. RVALUE reference is itself LVALUE, not RVALUE ! Use **std::move()** to move it!
7. The victim of MOVE is not destroyed, it becomes empty.

# Universal references

Universal ref can bind to both LVALUES and RVALUES:

```
template <typename T>
```

```
void fun(T && t){ ...}    // Universal REF !
```

```
auto && t = ...;    // Universal REF !
```

Only **T &&** or **auto &&** is universal ref, everything else with **&&** is RVALUE ref:

```
template <typename T>
```

```
void fun(vector<T> && t){ ...}    // Rvalue REF !
```

```
template <typename T>
```

```
void fun(T::type && t){ ...}    // Rvalue REF !
```

```
void fun(string && t){ ...}    // Rvalue REF !
```

# How does it work ???

Normally *REFERENCE to REFERENCE* is forbidden in C++:

```
int & & a = i; // ERROR !
```

But in template type deduction and **auto** the rules are different.

Reference collapse rules :

**& & -> &**

**&& & -> &**

**& && -> &**

**&& && -> &&**

In reality THERE IS NO UNIVERSAL REFERENCES.

They are special RVALUE references.

It is all about type deduction + reference collapsing.

# How does it work ???

Template with a universal ref :

```
template <typename T>  
void fun(T && t){ ...}    // Universal REF !
```

For LVALUE:

```
string s("Mickey Mouse");  
fun(s);
```

T is deduced as **string &**, and we get : **fun(string & && s);**

Which means **fun(string & s);**

For RVALUE:

```
string s("Mickey Mouse");  
fun(move(s));
```

T is deduced as **string** , and we get : **fun(string && s);**

# std::move possible implementation

In C++ 11:

```
template<typename T>
typename remove_reference<T>::type && move(T && param){
    using Return_Type = typename remove_reference<T>::type &&;
    return static_cast<Return_Type>(param); // Cast to RVALUE ref
}
```

In C++ 14:

```
template<typename T>
decltype(auto) move(T && param){
    using Return_Type = typename remove_reference<T>::type &&;
    return static_cast<Return_Type>(param); // Cast to RVALUE ref
}
```

But how does it work ?



# std::move for LVALUES

```
template<typename T>
typename remove_reference<T>::type && move(T && param){
    using ReturnTpe = typename remove_reference<T>::type &&;
    return static_cast<ReturnTpe>(param); // Cast to RVALUE ref
}
```

For string LVALUE: T = string & :

```
typename remove_reference<string &>::type && move(string & && param){
    using ReturnTpe = typename remove_reference<string &>::type &&;
    return static_cast<ReturnTpe>(param); // Cast to RVALUE ref
}
```

or

```
string && move(string & param){
    using ReturnTpe = string &&;
    return static_cast<ReturnTpe>(param); // Cast to RVALUE ref
}
```

# std::move for RVALUES

```
template<typename T>
typename remove_reference<T>::type && move(T && param){
    using ReturnType = typename remove_reference<T>::type &&;
    return static_cast<ReturnType>(param); // Cast to RVALUE ref
}
```

For string RVALUE: T = string :

```
typename remove_reference<string>::type && move(string && param){
    using ReturnType = typename remove_reference<string>::type &&;
    return static_cast<ReturnType>(param); // Cast to RVALUE ref
}
```

or

```
string && move(string && param){
    using ReturnType = string &&;
    return static_cast<ReturnType>(param); // Cast to RVALUE ref
}
```

# Perfect forwarding. Problem :

We want to write a template which adds element to a **vector** :

```
template<typename V, typename T>
void addToVec(V & v, const T & t){
    v.push_back(t);
}
```

Is it good ? Not exactly !

```
vector<Tjej> v;
v.reserve(10);
Tjej t1("Clair Lasbard");
addToVec(v, t1); // Add a LVALUE, 1 copy, OK
```

```
addToVec(v, Tjej("Nel Zelpher")); // RVALUE is copied instead of MOVE !!!
```

# Solution : Perfect forwarding + universal reference

*We perfect forward the argument :*

```
template<typename V, typename T>
void addToVec(V &v, T && t){
    v.push_back(forward<T>(t));
}
```

`std::forward()` always needs type argument : `forward<T>(...)`

*Is it good ? YES !*

```
vector<Tjej> v;
```

```
v.reserve(10);
```

```
Tjej t1("Clair Lasbard");
```

```
addToVec(v, t1); // Add a LVALUE, 1 copy, OK
```

```
addToVec(v, Tjej("Nel Zelpher")); // Add a RVALUE, 1 move, OK
```

# How does std::forward work ? Possible implementation.

```
template<typename T>
T && forward(typename remove_reference<T>::type & param){
    return static_cast<T &&>(param);
}
```

Type **T** is not deduced, but fixed by universal ref in e.g. **addToVec()**.

For string RVALUE: **T = string** : Similar to move :

```
string && forward(string & param){
    return static_cast<string &&>(param);
}
```

For string LVALUE: **T = string &** :

```
string & forward(string & param){
    return static_cast<string &>(param);
}
```

# Perfect forwarding variadic version

We want to write a template which adds element to a **vector** using **emplace\_back()**:

```
template <typename V, typename ... Args>
void addToVec2(V & v, Args && ... args){
    v.emplace_back(forward<Args>(args)...);
}
```

We perfect-forward all arguments to **emplace\_back()**.

```
addToVec2(v, string("Maria Traydor"));    // Construct in-place
```

## Linear algebra libraries in C++

- BLAS/LAPACK: classical Fortran libraries
- Eigen: header-only OR optionally can use BLAS/LAPACK for speed
- GNU Scientific Library (GSL): Copyleft, uses BLAS
- Armadillo: Nice C++ classes, uses BLAS/LAPACK
- OpenCV : Not really linear algebra, efficient linear algebra only with BLAS/LAPACK

Other languages:

- Python: numpy/scipy : uses BLAS/LAPACK
- Matlab : uses BLAS/LAPACK

Do you see any common theme here?

# Library of the day: CBLAS + LAPACK (or BLAS/LAPACK)

## **BLAS (Basic Linear Algebra Subprograms) :**

Very low-level operations (up to matrix\*matrix)

Original *Reference BLAS* : FORTRAN, slow

Modern BLAS is super-optimized for modern CPUs (SIMD etc.)

It is very important to use recent versions for modern CPUs !

*OpenBlas* : various architectures (Intel, AMD, ARM), open source

*Intel MKL* : Intel CPUs only, very slow on AMD, proprietary

## **LAPACK (Linear Algebra Package) :**

Higher-level operations (eigenvalues etc.)

Usually generic FORTRAN LAPACK is used

*Intel MKL* has an own modified LAPACK version

Fortran : BLAS + LAPACK

C/C++ wrappers : CBLAS+LAPACK (or you can use fortran libraries directly)




# BLAS+LAPACK

- Blas level 1: Vector-vector operations (vec+vec, vec.vec)
- Blas level 2: Matrix-vector operations (mat\*vec)
- Blas level 3: Matrix-matrix operations (mat\*mat)
- Lapack *computational routines*: QR-factorization etc.
- Lapack *driver routines*: Highest levels (eigenvalues+eigenvectors)


## BLAS+LAPACK facts:

- BLAS-LAPACK uses *column-major* matrices (Opposite to C/C++)!
- Support for special matrices (symmetric, band, triangular, sparse)
- Lapack has different algorithms to choose from for some operations
- Vectors and matrices can have strides ("increment" or "leading dimension")
- Uses low-level C (or Fortran), no C++ classes

Row-major order



Column-major order



Other libraries: FFTW (Fastest Fourier Transform in the West), ScaLapack (multi-CPU)

Documentation : <https://software.intel.com/en-us/mkl-developer-reference-c>

## Blas levels 1, 2 :

Blas level 1 example: vector-vector: dot product of two vectors (Note : C API, uses *pointers*)

```
vector<double> a{2., 3., 7., 2.}; // We can use vector
double b[] = {1., -1., 1., -2.}; // Or array
int n = 4; // Vector size
int inca = 1, incb = 1; // Strides of a, b
double result = cblas_ddot(n, a.data(), inca, b, incb); // Double DOT
```

Blas level 2 example: matrix-vector: operation **dgemv**  $\alpha * A * x + \beta * y$

**D**ouble **GE**neral matrix **M**atrix-**V**ector product

```
double a[]{1, 2, 3, 6, 5, 4};
double x[]{1, -1};
double y[]{2, 1, -1};
int m = 3, n = 2; // Dimensions of A
int incx = 1, incy = 1; // Strides of x, y
int lda = m; // "Leading dimensions" (matrix strides)
double alpha = -1., beta = -2.;
cblas_dgemv(CblasColMajor, CblasNoTrans, m, n, alpha,
            a, lda, x, incx, beta, y, incy);
```

## Blas levels 3 :

Blas level 3 example: matrix-matrix: operation **dgemm** (  $\alpha * A * B + \beta * C$  )

**D**ouble **G**eneral matrix **M**atrix-**M**atrix product

```
double a[]{1, 2, 3, 6, 5, 4}; // 3x2
double b[]{1, -1, 2, 0, 0, 2, 1, 1}; // 2x4
double c[]{0.5, 0.5, 0.5, 0.5, 0.5, 0.5, 0.5, 0.5, 0.5, 0.5, 0.5, 0.5, 0.5}; // 3x4

int m = 3, k = 2; // Dimensions of A
int n = 4; // Columns of B
int lda = m, ldb = k, ldc = m; // "Leading dimensions" (matrix strides)
double alpha = 0.5, beta = 2.0;

cblas_dgemm(CblasColMajor, CblasNoTrans, CblasNoTrans, m, n, k,
            alpha, a, lda, b, ldb, beta, c, ldc);
```

# Lapack

Lapack example: matrix-matrix: diagonalize a real symmetric matrix : operation **dsyevr**

**D**ouble **S**Ymmetric **E**igen**V**alues with Relatively **R**obust Representations

This is a high-level driver routine!

```
double a[2][2] = {0., 1., 1, 0.};
int n = 2; // Dimension of a
int lda = 2; // Leading dimension of z
char jobz = 'V'; // Compute e-values + e-vectors
char range = 'A'; // Compute all e-values
char uplo = 'U'; // Upper triangular part of a

double z[2][2]; // Output eigenvectors
int ldz = 2; // Leading dimension of z
double w[2] {-3.14, -3.14}; // Output eigenvalues
int m = -1; // Output: number of e-values found
int isuppz[4]; // "Support output"

LAPACKE_dsyevr(LAPACK_COL_MAJOR, jobz, range, uplo, n, (double *)a, lda, 0, 0,
               0, 0, 0, &m, w, (double *)z, ldz, isuppz);
```

**title**

text

**Thank you for your attention !**

**THE END**